

Azure Open Source Day

Improve the development for
.NET application & Azure Cosmos DB for PostgreSQL



Who I am



Giorgio Desideri

Tech Lead of Cloud Solution

Seven Peaks Software



Agenda


Table of contents

1. Working with PostGreSQL Cluster
2. Development of the application
3. Application Use-cases



Working with a PostgreSQL Cluster

Who ever work
with PGSQL
Cluster or any
other SQL-DB
Cluster ?



Working with a PostGreSQL Cluster

Issues

- Understand the concept of “**Shard**” or “**Partitioning**”
- Performances and how to obtain them
- Monitoring of what the application is doing
 - and how the application does it



Working with a PostgreSQL Cluster

Issues

- Understand the concept of “**Shard**” or “**Partitioning**”
- Performances to retrieve and write data
 - ... and how to obtain them
- Monitoring of what the application is doing
 - ... and how the application does it

Do you have any other ?



Development of the application

Head scratch points

- Domain / Entities design
- Data Access
- Performances related to the Domain / Entities operations



“Head Scratch” = *think hard in order to find a solution to something.*



Entities design

Reasons

- Cosmos DB PostgreSQL has 3 different types of tables:
 - Distributed tables
 - Reference tables
 - Local tables



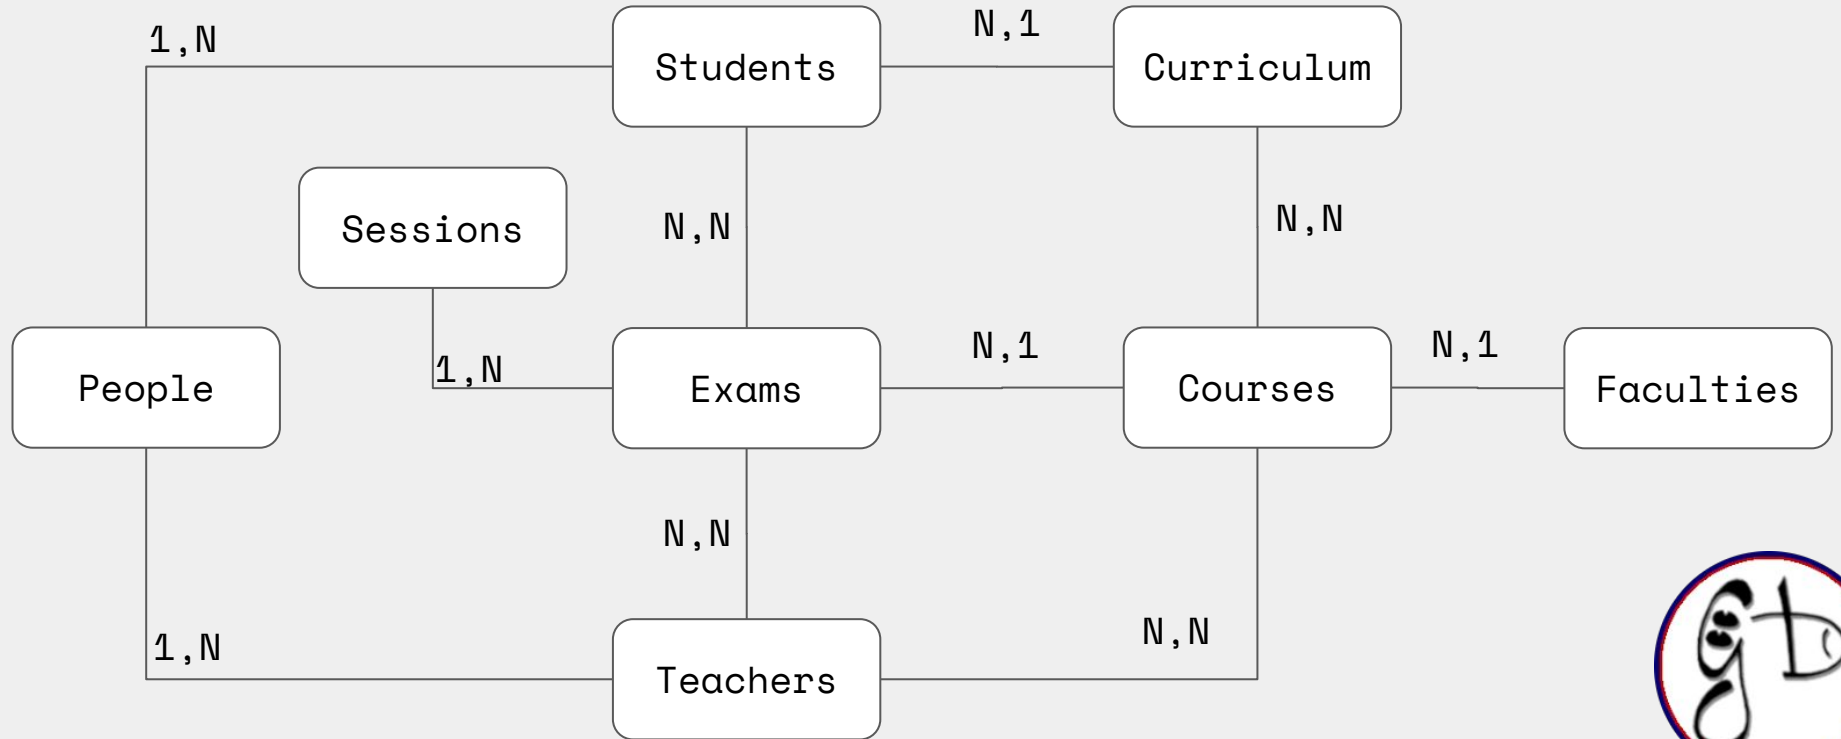
Entities design

Challenge

- What (entities) can be “*distributed*” ?
- What (entities) can be “*reference*” ?
- What (entities) can be “*local*” ?



Entities design



Entities design

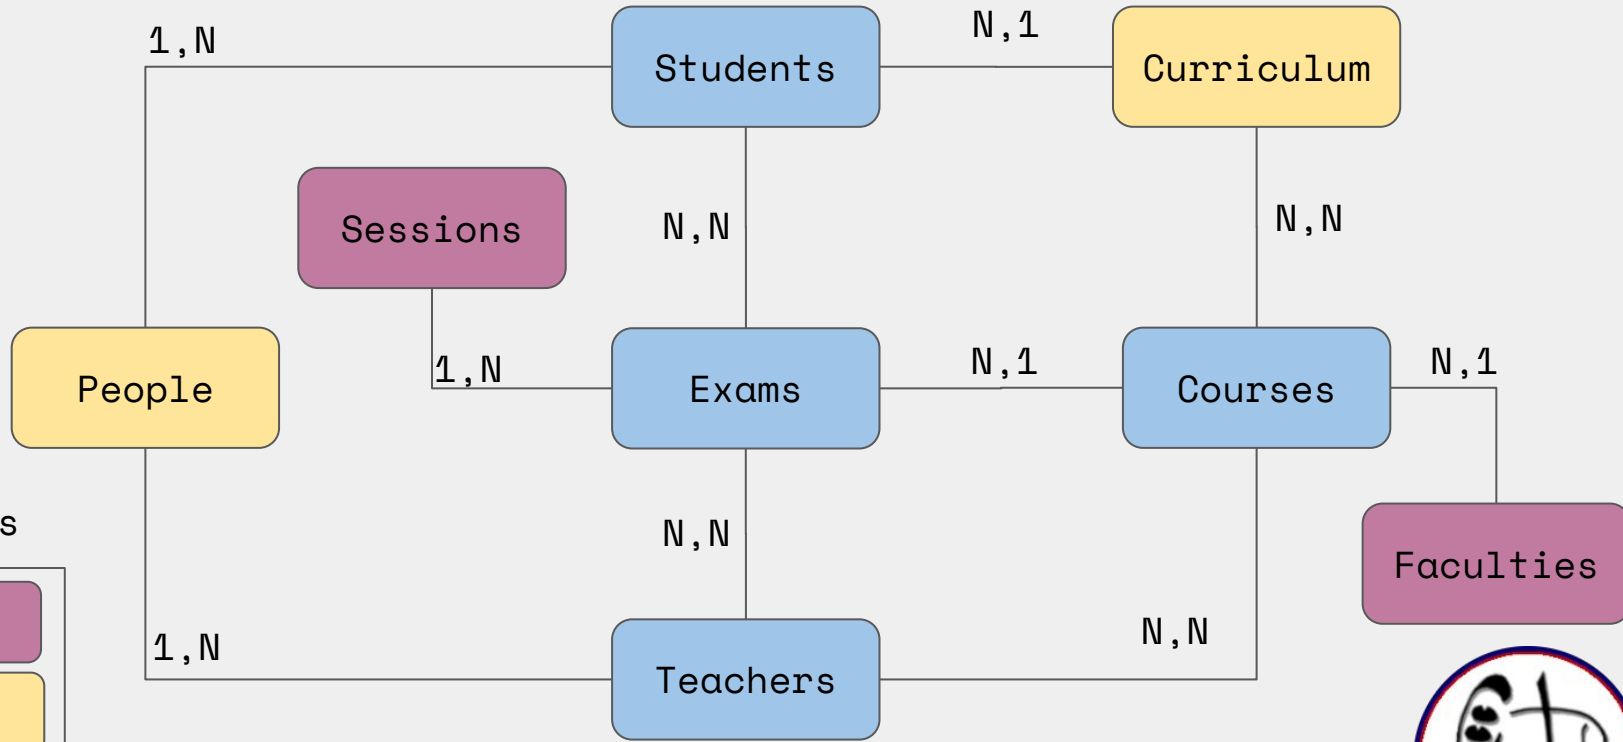


Table types

Reference

Local

Distributed



Entities design

Key concepts

- Understand the point of view for distribution
 - Sharding
 - Table colocation
 - Performance for read and write
 - Costs of Time complexity
 - Amount of Nodes
 - Total RUs



Entities design

Key Concepts

- SQL Normalization theory (1NF,2NF,3NF,...)
 - Difference between reference & not-reference entities
 - Using Entities \Rightarrow No good choice for cluster, because move semantic logic from data layer to business layer.
 - Cost of Space complexity
 - RUs used to retrieve or write
 - Table dimension (space allocation)



Entities design

Key Concepts

- Hierarchy analysis
 - TPH : Table-per-Hierarchy
 - TPT : Table-per-Type
 - TPC : Table-per-Concrete Type
 - Costs over space and time complexity



Data Access

Reason

- SQL vs ORM
- Connection Pool



Data Access

Challenge

- Usage of SQL brings:
 - More control over the queries
 - More flexibility over the DDL
 - Native support of the cluster commands
- Usage of ORM brings:
 - Code-first approach
 - Configuration over the ORM interfaces
 - Supervisioning over the migrations



Data Access

Challenge

- Connection Pools
 - It is a cluster, not a database !
 - Differences of implementation are performance-driven
 - Connection reuse
 - Fast connection acquisition
 - *Open()* or *Close()* automatic management



Data Access

Key Points

- All is performance-driven, so considering to study and analyse your queries
- SQL approach is the most “balanced” way
 - to use something that you “know”
 - reduce the risk to spend X amount of hours in customisation / overriding / work-arounding



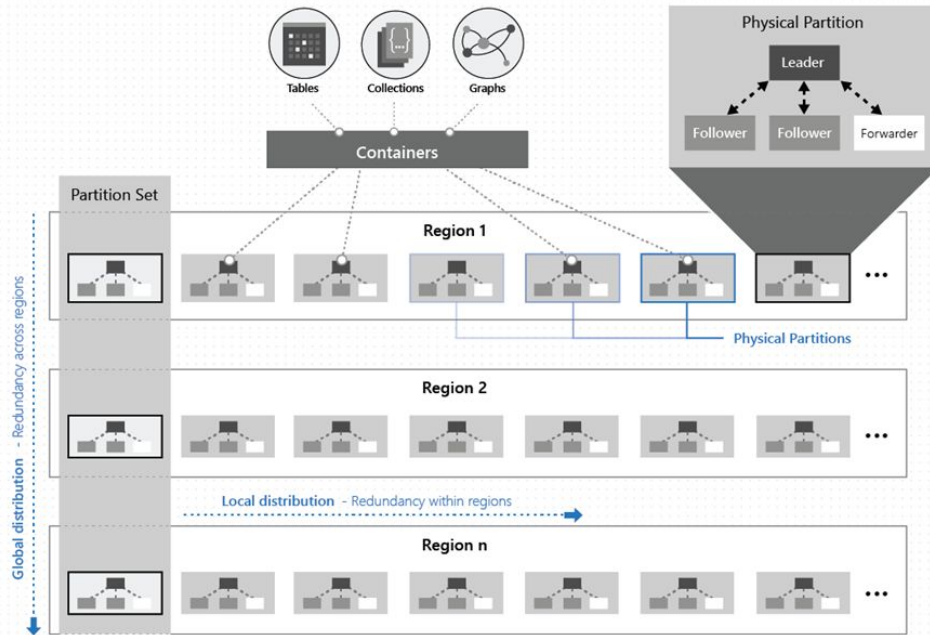
Application Use-Cases

Reasons

- Multi-tenant SaaS
 - => Shard key identification
 - => RU understanding
- Real-Time application
 - => OLTP (On-Line Transaction Processing)
 - => Concurrency



Multi-Tenant SaaS



Shard Key Identification

- Logical partition
- Physical partition
- Partition Set
- Replication
- Consistency Level



Multi-Tenant SaaS

RUs understanding

- RU = Request Unit, $1 \text{ RU} \approx 1 \text{ Kb/s}$
- RU has impact over
 - all operations (Read / Write)
 - data model
- PoC-driven approach for development
 - Cannot be estimated OR Cannot be forecast
 - Must be implemented
- Concept of No-SQL DB usage



Real-Time application

Processing models

- OLTP : On-Line Transaction Processing
 - execute a number of transactions occurring concurrently
 - Transaction Lock
- OLAP : On-Line Analytical Processing
 - performing multidimensional analysis at high speeds on large volumes of data from a data warehouse, data mart
 - Roll-up / Drill-Down / Slice / Dice / Pivot




Real-Time application

Usage

Citus brings the power of **distributed tables** to PostgreSQL

APPLICATION

```
CREATE TABLE campaigns (...);  
SELECT create_distributed_table(  
    'campaigns', 'company_id');
```



WORKER NODES

CREATE TABLE
campaigns_101
CREATE TABLE
campaigns_104



CREATE TABLE
campaigns_102
CREATE TABLE
campaigns_105



CREATE TABLE
campaigns_103
CREATE TABLE
campaigns_106



Links

- <https://learn.microsoft.com/en-us/azure/cosmos-db/postgresql/>
- <https://learn.microsoft.com/en-us/azure/cosmos-db/postgresql/concepts-cluster#nodes>
- <https://learn.microsoft.com/en-us/azure/cosmos-db/postgresql/concepts-nodes>
- <https://learn.microsoft.com/en-us/azure/cosmos-db/postgresql/quickstart-build-scalable-apps-overview>
- <https://learn.microsoft.com/en-us/azure/cosmos-db/postgresql/quickstart-build-scalable-apps-concepts>
- <https://learn.microsoft.com/en-us/azure/cosmos-db/postgresql/quickstart-build-scalable-apps-model-multi-tenant>
- <https://learn.microsoft.com/en-us/azure/cosmos-db/postgresql/resources-pricing>



The End

Q / A

